

# STAT 4830: Numerical optimization for data science and ML

## Lecture 4: Beyond Least Squares

From Manual to Automatic Differentiation

Professor Damek Davis

# The Problem: Manual Gradient Computation

Consider computing this gradient by hand:

$$f(w) = \frac{1}{2} \| \tanh(W_2 \text{ReLU}(W_1 x + b_1) + b_2) - y \|^2$$

Challenges:

- Complex chain rule applications
- Error-prone derivations
- Time-consuming process
- Limited to simple functions

# The Solution: Automatic Differentiation

PyTorch provides:

```
# Define complex function
def f(x, W1, b1, W2, b2):
    h = torch.relu(W1 @ x + b1)
    return 0.5 * torch.sum(
        (torch.tanh(W2 @ h + b2) - y)**2
    )

# Get gradient automatically
f.backward()
```

Key benefits:

1. Automatic gradient computation
2. Handles any differentiable function
3. Memory efficient implementation
4. Scales to large problems

# Three Key Ideas

1. **Computational Graph**
2. **Reverse-Mode Differentiation**
3. **Memory-Efficient Implementation**

# Outline

## 1. Computing Gradients

Function → Graph → Gradient

## 2. Gradient Descent

Gradient → Update → Repeat

## 3. Neural Networks

Features → Layers → Loss

# A Simple Example: Polynomial Function

Let's start with a one-dimensional function:


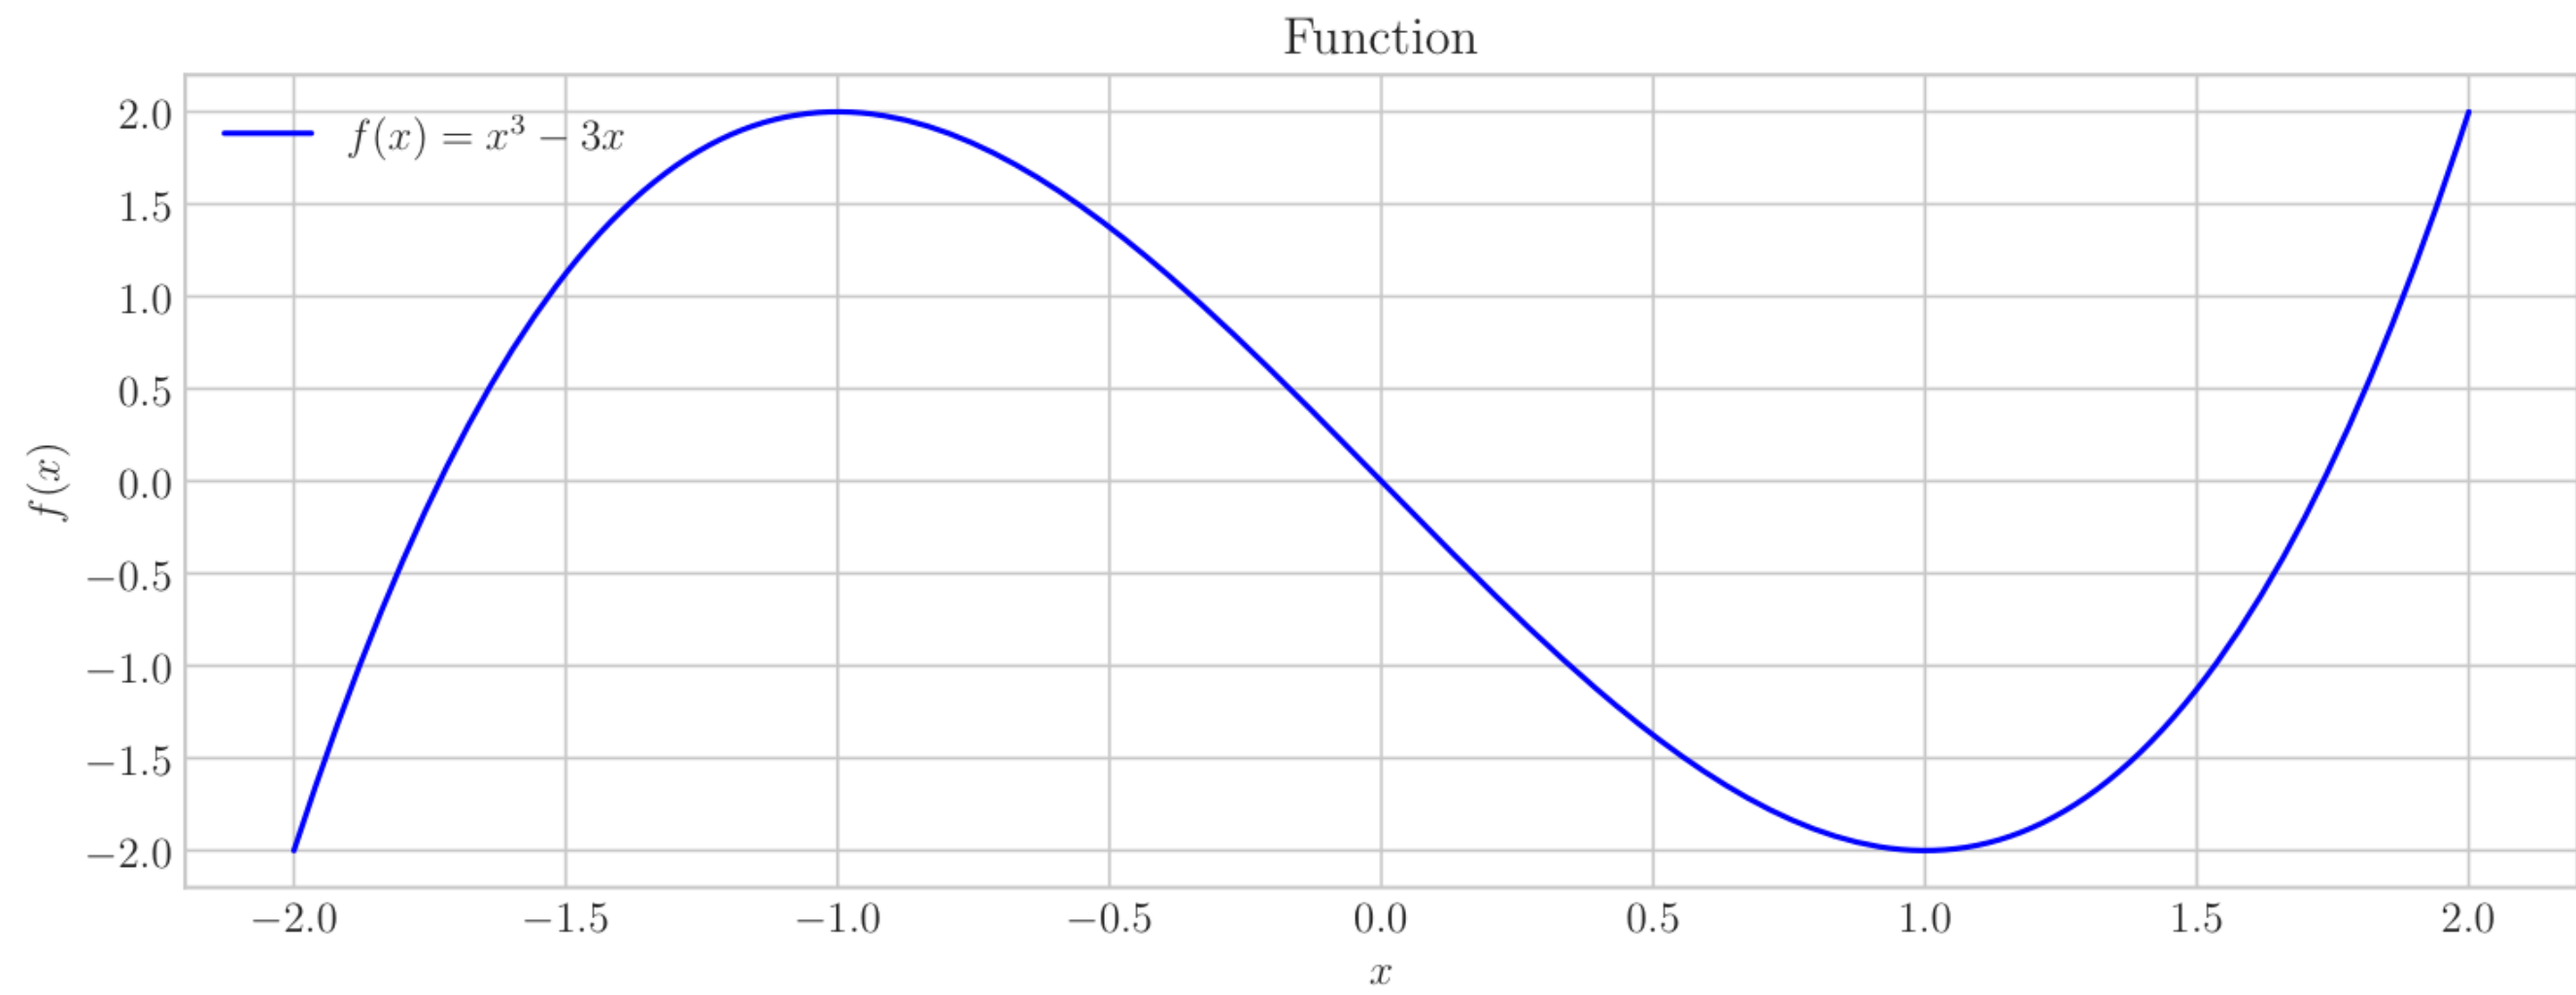
$$f(x) = x^3 - 3x$$

Manual gradient computation:

$$\frac{d}{dx} f(x) = 3x^2 - 3$$

PyTorch automates this:

```
x = torch.tensor([1.0], requires_grad=True) # Gradient Tracking
y = x**3 - 3*x # Forward Pass
y.backward() # Backward Pass
print(f"f'(1) = {x.grad}") # Gradient Access
```

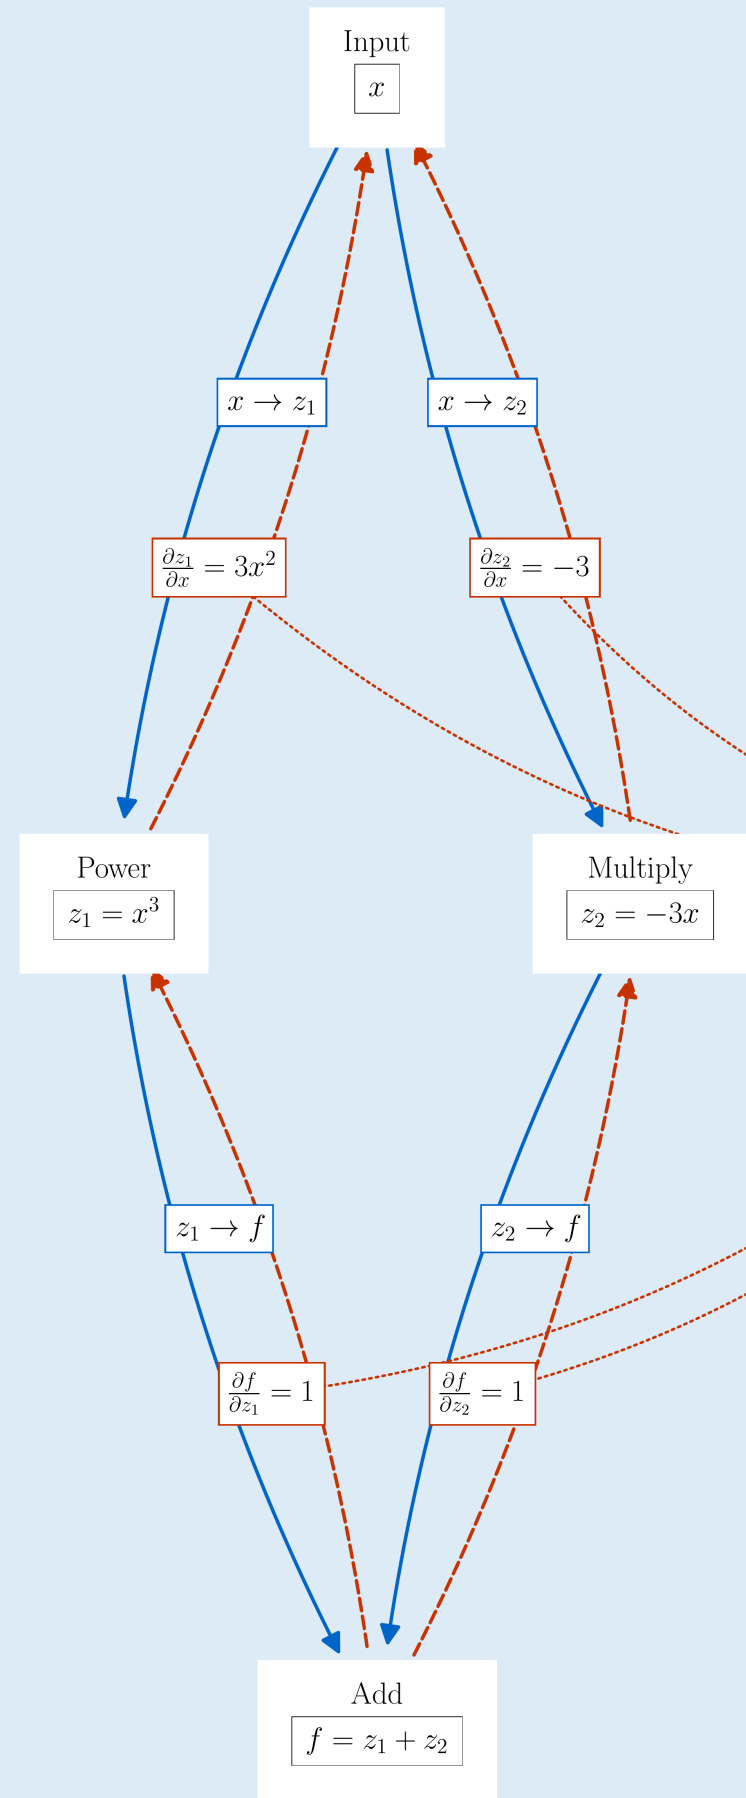


# How does PyTorch do this?


- **Forward pass:** When you evaluate a function, PyTorch computes a *computational graph* that records all operations like addition, multiplication, powers, etc.
- **Backward pass:** PyTorch traverses the graph in reverse order to compute the gradient, using what is essentially an efficient implementation of the chain rule.



Forward Pass  
(builds dynamic computation graph)



Backward Pass  
(computes gradients via chain rule)



# Building the Computational Graph

Each node in the graph:

- Stores output value from forward pass
- Contains function for local gradients
- Maintains references to inputs

The graph structure:

- Records operations
- Stores values
- Enables gradient flow

For  $f(x) = x^3 - 3x$ , we build:

1. Input node storing  $x$
2. Power node computing  $z_1 = x^3$
3. Multiply node computing  $z_2 = -3x$
4. Add node forming  $f = z_1 + z_2$

# Computing Gradients: The Process

## Starting State:

- Initialize  $\frac{\partial f}{\partial f} = 1$  at output
- All other gradients start at 0

## Algorithm:

1. Process nodes in reverse order
2. Compute local gradients
3. Multiply by incoming gradient
4. Add to input gradients

## Key Features:

- Reverse topological sort
- Chain rule at each step
- Gradient accumulation
- Memory efficient

**Output Node** ( $f = z_1 + z_2$ ):

# Gradient Flow: Step by Step

- $\frac{\partial f}{\partial f} = 1$
- $\frac{\partial f}{\partial z_1} = 1, \frac{\partial f}{\partial z_2} = 1$
- Propagate to both input nodes

**Power Node** ( $z_1 = x^3$ ):

- Incoming gradient: 1
- Local gradient:  $\frac{\partial z_1}{\partial x} = 3x^2$
- Contribute:  $\frac{\partial f}{\partial x} += 3x^2$

**Multiply Node** ( $z_2 = -3x$ ):

- Incoming gradient: 1
- Local gradient:  $\frac{\partial z_2}{\partial x} = -3$
- Contribute:  $\frac{\partial f}{\partial x} += -3$

**Input Node** ( $x$ ):

# Building a Computational Graph

Let's see how PyTorch builds a graph for:

$$f(x) = x^3 - 3x$$

Step 1: Create input node

```
x = torch.tensor([1.0],  
                  requires_grad=True)
```

Key properties:

- Tracks gradients
- Stores value
- Records operations

# Building a Computational Graph

Step 2: Power operation  $z_1 = x^3$

```
z1 = x**3
```

Graph grows:

- New operation node
- Stores intermediate value
- Records connection to input

# Building a Computational Graph

Step 3: Linear term  $z_2 = -3x$

```
z2 = -3 * x
```

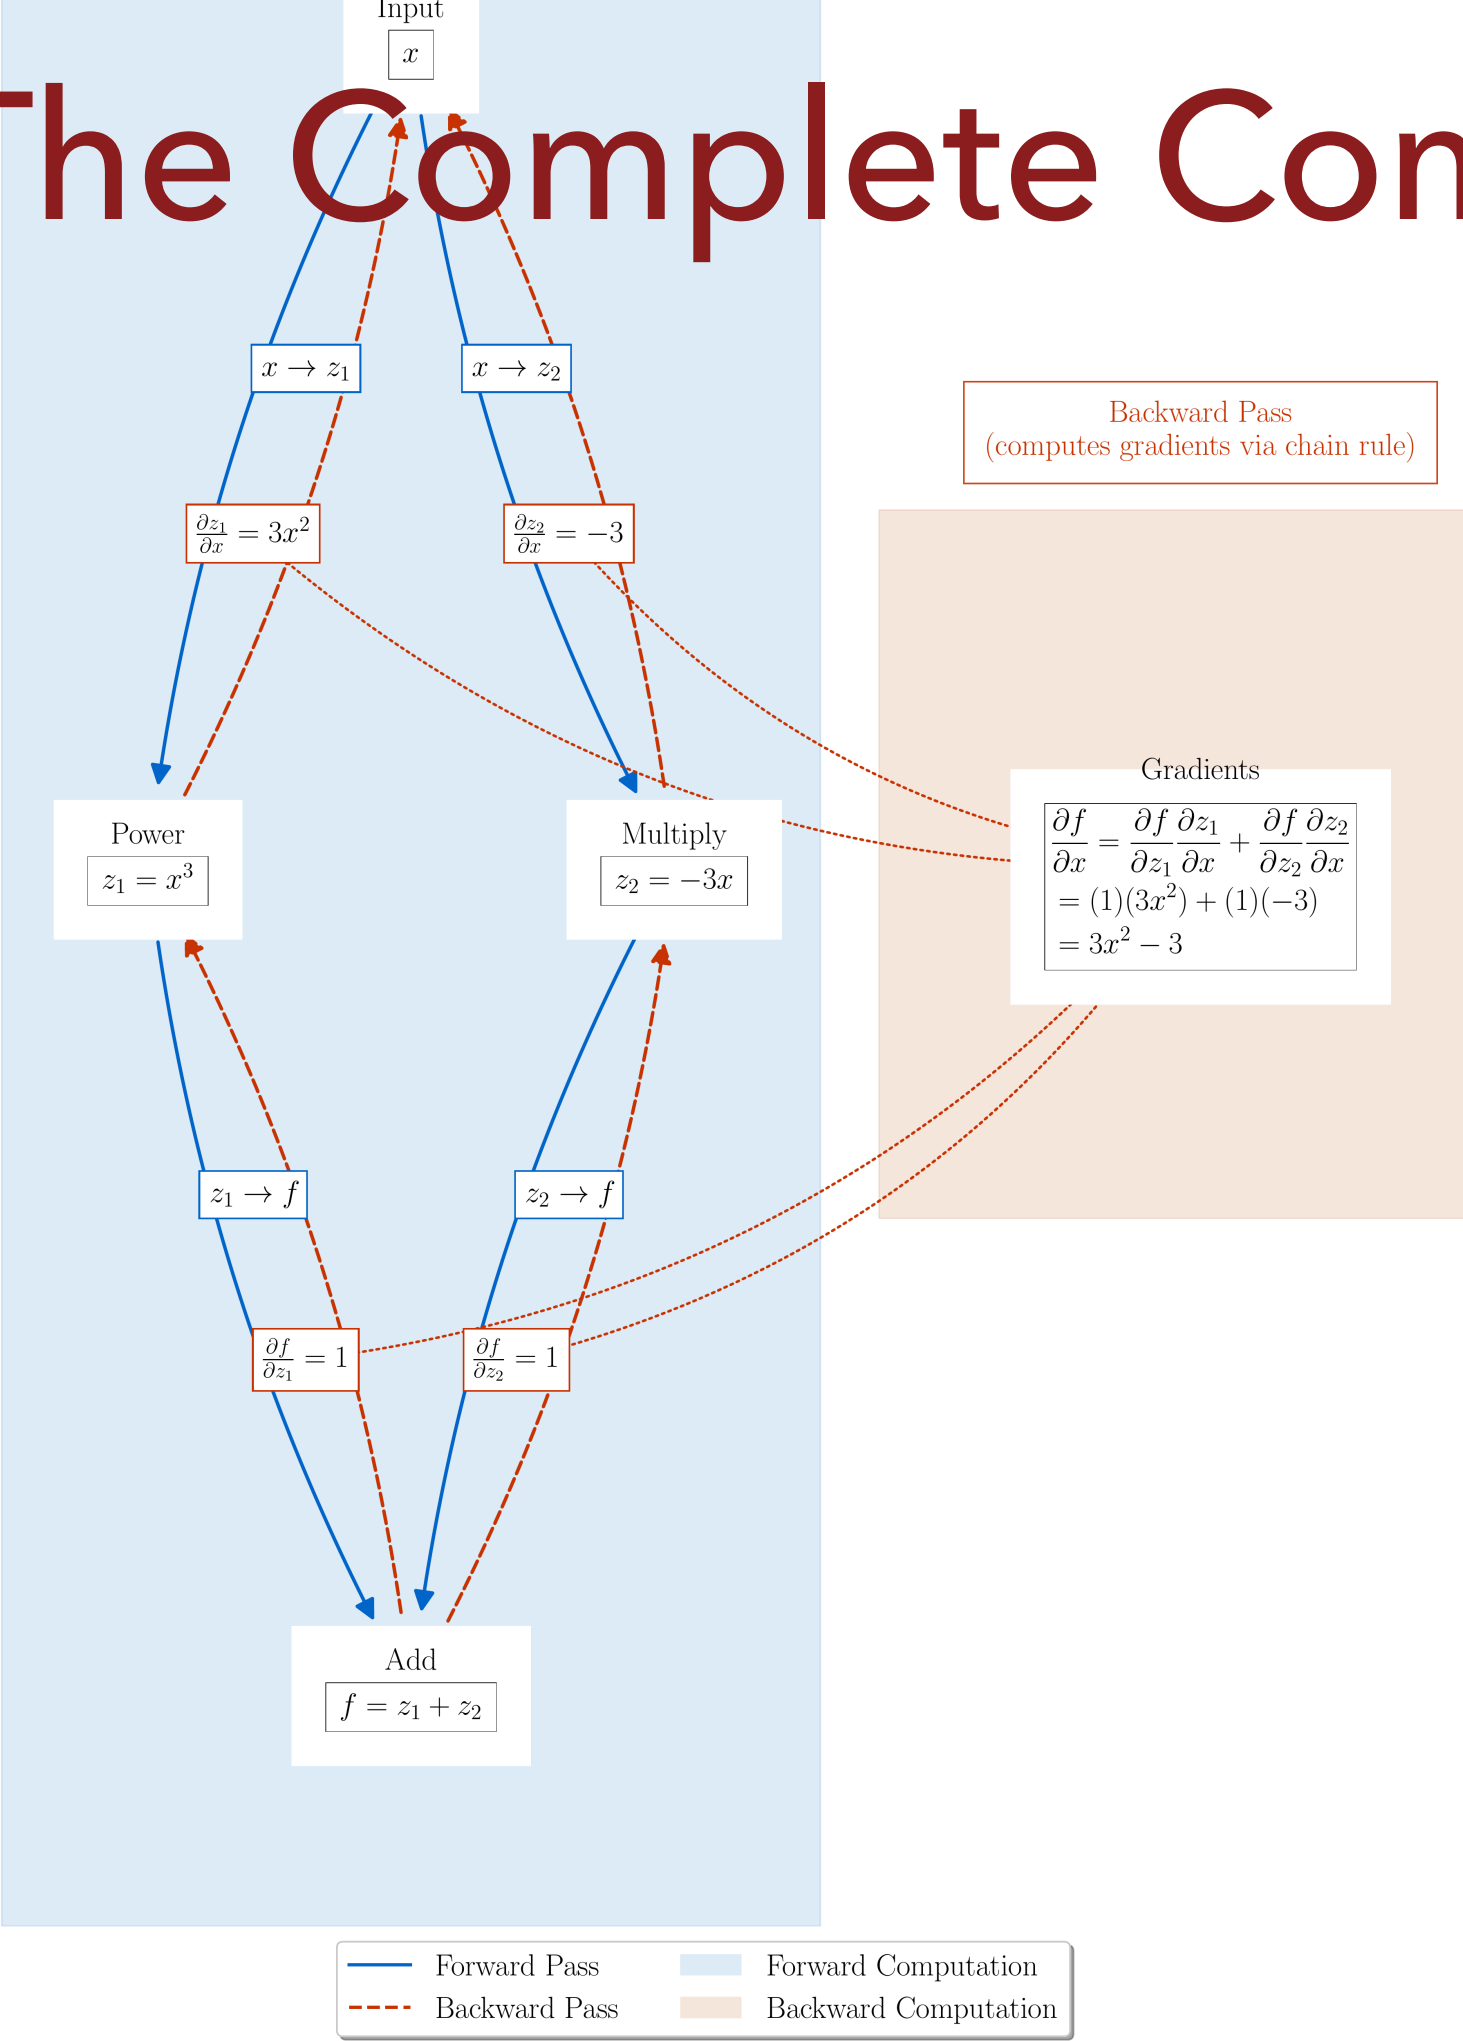
Note:

- Reuses input node
- Creates new operation
- Stores scalar multiplier

Multiple paths from x:

- Through power node
- Through multiply node
- Will need accumulation

# The Complete Computational Graph





For input  $x = 1$ :

# Forward Pass: Computing Values

1.

Power node:

$$z_1 = 1^3 = 1$$

2.

Multiply node:

$$z_2 = -3(1) = -3$$

3.

Add node:

$$f = 1 + (-3) = -2$$

*# Forward computation*

```
x = torch.tensor([1.0],  
                  requires_grad=True)
```

```
z1 = x**3
```

```
z2 = -3 * x
```

```
f = z1 + z2
```

```
print(f"z1: {z1.item()}") # 1.0
```

```
print(f"z2: {z2.item()}") # -3.0
```

```
print(f"f: {f.item()}") # -2.0
```

# Backward Pass: Computing Gradients

Starting from output:

1. Initialize  $\frac{\partial f}{\partial f} = 1$

2. Flow through power path:

$$\frac{\partial f}{\partial x} += 3x^2$$

3. Flow through multiply path:

$$\frac{\partial f}{\partial x} += -3$$

Gradient accumulation:

- Sum contributions
- Multiple paths
- Chain rule at each step

# The Chain Rule: A Visual Guide

For a chain of operations:

$$x \xrightarrow{g} z \xrightarrow{h} y$$

The chain rule states:

$$\frac{dy}{dx} = \frac{dy}{dz} \cdot \frac{dz}{dx}$$

Gradient flows backward:

1. Start at output
2. Multiply derivatives
3. Follow paths back

# Chain Rule in PyTorch

```
# Forward pass  
z = g(x)  # First function  
y = h(z)  # Second function  
  
# Backward pass  
dy_dz = h.grad_fn(z)    # Local grad  
dz_dx = g.grad_fn(x)    # Local grad  
dy_dx = dy_dz * dz_dx    # Chain rule
```

Each node stores:

1. Forward function
2. Gradient function
3. Input references

PyTorch handles:

- Function composition
- Gradient computation
- Memory management

# Two Implementation Approaches

## 1. Using `backward()`

```
# Create graph
x.requires_grad = True
z = g(x)
y = h(z)

# Compute gradients
y.backward()
grad = x.grad # Stored in tensor
```

Best for:

- Training loops
- Multiple gradients
- Memory efficiency

## 2. Using `autograd.grad()`

```
# Create graph
x.requires_grad = True
z = g(x)
y = h(z)

# Direct computation
grad = torch.autograd.grad(y, x)[0]
```

Best for:

- One-off gradients
- Direct access
- Higher derivatives

# Memory Management: Theory vs Practice

Manual gradient:

```
# Forms huge matrices  
XtX = X.T @ X #  $O(p^2)$  memory  
grad = XtX @ w # Matrix-vector
```

Problems:

- Excessive memory use
- Poor cache utilization
- Limited scalability

PyTorch gradient:

```
# Matrix-vector only  
Xw = X @ w #  $O(p)$  memory  
grad = X.T @ Xw # Matrix-vector
```

Benefits:

- Minimal memory use
- Cache-friendly
- Scales to large problems

# During Training Best Practices for Memory

```
optimizer.zero_grad()  # Clear gradients  
loss.backward()        # Compute gradients  
optimizer.step()       # Update weights
```

2.

## During Evaluation

```
with torch.no_grad():  # No gradients needed  
    model.eval()       # Evaluation mode  
    predictions = model(data)
```

3.

## Memory Management

# From Simple to Complex: Least Squares

Manual gradient:

$$\nabla f = X^{\top} (Xw - y)$$

Requires:

- Matrix formation
- Careful derivation
- Memory allocation

PyTorch gradient:

```
pred = X @ w
loss = 0.5*((pred - y)**2).sum()
loss.backward()
grad = w.grad
```

Benefits:

- Automatic computation
- Memory efficient
- Scales naturally



# The Least Squares Graph: Step by Step

2.

Subtract:

$$\mathbf{z}_2 = \mathbf{z}_1 - \mathbf{y}$$

- Input:  $\mathbf{z}_1, \mathbf{y} \in \mathbb{R}^n$

- Output:  $\mathbf{z}_2 \in \mathbb{R}^n$

3.

# Gradient Flow in Least Squares

Total derivatives:

1.  $\frac{\partial f}{\partial z_3} = \frac{1}{2}$
2.  $\frac{\partial z_3}{\partial \mathbf{z}_2} = 2\mathbf{z}_2^\top$
3.  $\frac{\partial \mathbf{z}_2}{\partial \mathbf{w}} = \mathbf{X}$

Chain rule:

$$\frac{\partial f}{\partial \mathbf{w}} = \frac{1}{2} \cdot 2\mathbf{z}_2^\top \cdot \mathbf{X}$$

Key insights:

1. Row vector gradients
2. Matrix-vector products
3. No matrix formation
4. Memory efficient

Final gradient:

$$\nabla f = \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$

# Building Neural Networks: The Architecture

Layer composition:

Input  $\rightarrow$  Linear<sub>1</sub>  $\rightarrow$  Tanh  $\rightarrow$  Linear<sub>2</sub>  $\rightarrow$  Output  
 $\mathbb{R}^d$   $\mathbb{R}^{h \times d}$   $\mathbb{R}^h$   $\mathbb{R}^{1 \times h}$   $[0, 1]$

Each layer:

1. Linear transform
2. Nonlinear activation
3. Gradient tracking

PyTorch handles:

- Parameter management
- Forward computation
- Backward gradients

# Neural Network Implementation

```
def forward(self, x):  
    # Hidden features  
    h = torch.tanh(self.linear1(x))  
    # Output probability  
    return torch.sigmoid(  
        self.linear2(h)  
    )
```

- Memory optimization


2.

## Layer Organization

- Modular design
- Easy composition
- Clear data flow

3.

# MNIST Classification: The Task



## Dataset:

- 60,000 training images
- 10,000 test images
- 28×28 pixels each
- Binary labels (odd/even)

## Preprocessing:

```
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(
        (0.1307,), (0.3081,)
    )
])

# Load data
train_dataset = datasets.MNIST(
    './data',
    train=True,
    transform=transform
)
```

# Model Comparison: Architecture

## Logistic Regression:

```
class Logistic(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.linear = nn.Linear(784, 1)  
  
    def forward(self, x):  
        # Single linear layer  
        return torch.sigmoid(  
            self.linear(x.view(-1, 784))  
        )
```

## Neural Network:


```
class SimpleNN(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.fc1 = nn.Linear(784, 32)  
        self.fc2 = nn.Linear(32, 1)  
  
    def forward(self, x):  
        # Hidden layer with ReLU  
        h = torch.relu(  
            self.fc1(x.view(-1, 784))  
        )  
        # Output layer  
        return torch.sigmoid(self.fc2(h))
```

# Training Process: Step by Step

```
def train_model(model, train_loader, optimizer, epochs=5):  
    criterion = nn.BCELoss()  
    for epoch in range(epochs):  
        for batch_idx, (data, target) in enumerate(train_loader):  
            # 1. Zero gradients  
            optimizer.zero_grad()  
  
            # 2. Forward pass  
            output = model(data)  
            loss = criterion(output, target.float())  
  
            # 3. Backward pass  
            loss.backward()  
  
            # 4. Update weights  
            optimizer.step()  
  
            # 5. Log progress
```

PyTorch handles all gradient computation automatically.

# Results Analysis



## Final Results:

- Logistic: 87.30% accuracy
- Neural Net: 92.40% accuracy

## Key differences:

1. Feature learning
2. Nonlinear boundary
3. Better capacity

## Training progress:

- Faster neural net learning
- Higher final accuracy
- Better generalization



# Key Takeaways

- Scales to complex networks

2.

## Memory Efficiency

- Never forms large matrices
- Uses matrix-vector products
- Enables large-scale optimization
- Scales to deep networks

# Next Steps

- Stochastic gradients
- Adaptive methods
- Second-order techniques

2.

## Deep Learning

- Complex architectures
- Custom loss functions
- Training strategies

# Questions?

- Course website: <https://damek.github.io/STAT-4830/>
- Office hours: Listed on course website
- Email: [damek@wharton.upenn.edu](mailto:damek@wharton.upenn.edu)
- Discord: Check email for invite